

iRDRC: Markov Decision Process Deep Dive & Implementation

BTP Discussion 2

Joel James & Majida Bee

29 August 2026

Focus: Complete mathematical formulation of the MDP, why it works, and how to implement it.

Recap: The iRDRC Problem

Aspect	Description
Problem	AV needs radar + communication on shared resource
Environment	Weather, road, speed, moving objects (binary states)
Objective	Maximize throughput while minimizing miss-detection risk
Approach	Adaptive learned policy via Deep Q-Network (DQN)
Key Insight	Fixed scheduling fails under dynamic uncertainty

Meeting 1 Covered: System model (queue, channel, radar), environment model (Eq. 1), and introduced MDP components. **Today's Focus:** Deep understanding of MDP formulation, why each component matters, and how DQN solves it.

Why Markov Decision Process (MDP)?

The Markov Property: The future depends only on the current state, not on history.

Why this matters for iRDRC:

1. Queue state d captures the current backlog; only d matters for next queue size, not when packets arrived
2. Channel state c is independent each step; history doesn't help predict next c
3. Environmental factors (r, w, v, m) are resampled; each time step is a fresh draw
4. Radar detection is perfect; if radar is on and event exists, we catch it regardless of past misses

Result: Given current state s and action a , the next state s' is probabilistically determined, allowing us to use RL algorithms that exploit the Markov property.

The MDP Tuple: ■S, A, P, R, γ ■

A Markov Decision Process is defined by five components. Together, they describe a complete sequential decision problem:

Component	Symbol	Meaning (in iRDRC)	Example
State Space	S	All possible observations the AV can make	$\{c=5, m=0, r=0, w=1, v=1, m=0\}$
Action Space	A	All possible choices the AV can make	$\{0: no\ communication, 1: radar\}$
Transition Prob.	$P(s' s,a)$	How state evolves given action	$P(d=4 \mid d=5, a=0, c=0)$
Reward Function	$R(s,a,s')$	Immediate reward for (state, action, next state)	$50 - n_e \times (goal - comm), -50$ (missed event)
Discount Factor	γ	Weight on future vs. immediate reward	0.99

State Space (Equation 2)

Definition: $S = \{ (d, c, r, w, v, m) : d \in \{0, \dots, 10\}, c, r, w, v, m \in \{0, 1\} \}$

State Space Size Calculation:

- d (queue): 11 possible values (0, 1, 2, ..., 10)
- c (channel): 2 possible values (0 = good, 1 = bad)
- r (road): 2 possible values
- w (weather): 2 possible values
- v (speed): 2 possible values
- m (moving objects): 2 possible values

Total: $11 \times 2^5 = 352$ states

Key Insight: 352 states is manageable but large enough to benefit from function approximation (i.e., neural networks) instead of tabular methods.

Action Space (Simple but Powerful)

Definition: $A = \{0, 1\}$

What Each Action Means:

Action	Mode	What Happens	Performance Impact
$a = 0$	Communication	Transmit queued packets; channel detection is skipped	Increases throughput, increases miss-detection risk
$a = 1$	Radar	Scan for obstacles; if event exists and radar is active, transmit	Decreases throughput, decreases miss-detection risk

The Central Design Decision: Each time step, the AV chooses exactly one mode. This binary action is the lever that the learning algorithm will adjust to find the optimal balance.

Reward Function (Equation 3) — The Learning Signal

The reward function translates engineering objectives (throughput + safety) into numbers that a learning algorithm can optimize. It is the bridge between goals and learning.

Action	Situation	Reward	Interpretation
Comm (a=0)	Good channel, no event	$+r_{\text{good}} = +2$	Success: 4 packets sent
Comm (a=0)	Bad channel, no event	$+r_{\text{bad}} = +1$	Partial: 2 packets sent
Comm (a=0)	Unexpected event occurs	$-r_{\text{event}} = -50$	Catastrophic: missed obstacle
Radar (a=1)	No event occurs	0	Radar cost: detection found nothing
Radar (a=1)	Event occurs (b unfav.)	$+r_{\text{event}}(b+1) = +5(b+1)$	Success: detected threat

reward scales with risk

Key Design Choices:

- **-50 r_{event} , r_{good} (Safety-First):** Missing an event is 25–50× worse than missing a few packets. This asymmetry encodes a safety priority.
- **Reward scales with risk (b):** A detected event in high-risk (b=4) earns +25, but in low-risk (b=0) earns only +5. Encourages the policy to patrol risky conditions.
- **Radar with no event = 0:** No reward, but no penalty either. Radar is insurance against uncertainty.

Objective: Discounted Return (Equation 4)

Goal: Find a policy π that maximizes the expected cumulative reward over many time steps.

Discounted Return (Value of a Policy π):

$$G(\pi) = E\left\{ \sum_{t=0}^T \gamma^t r_{t+1}(\pi) \right\}$$

Components:

- $r_{t+1}(\pi)$: Immediate reward at time $t+1$ under policy π
- $\gamma \in (0,1)$: Discount factor (paper doesn't specify; typical ≈ 0.99)
- γ^t : Exponential decay; rewards 10 steps away are worth $\gamma^{10} \approx 0.90$ compared to immediate
- T : Time horizon (episode length; likely 500–1000 steps)

Optimal Policy: $\pi^* = \arg \max_{\pi} G(\pi)$

The policy that, when followed, yields the highest expected cumulative reward.

Intuition: Don't chase immediate rewards. Choose actions that lead to long-term payoff. If communication now leads to a crash later, the -50 penalty far outweighs the $+1$ or $+2$ gain.

Why Discount the Future? (γ Explained)

What γ Controls:

- $\gamma \rightarrow 1.0$ (e.g., 0.99): Future rewards nearly as valuable as immediate rewards. Policy is far-sighted. Over time horizon $T=1000$, reward at step 999 is still worth $\gamma^{1000} \approx 0.0005$ of immediate reward.
- $\gamma \rightarrow 0.0$ (e.g., 0.5): Only immediate rewards matter. Policy is myopic; will choose any immediate gain even if it ruins future prospects.

For iRDRC:

- $\gamma \approx 0.99$ makes sense because a crash today has consequences for all future steps.
- If γ were 0.5, the policy might communicate recklessly now to gain +2, knowing the -50 crash penalty is only 50% as important.
- $\gamma \approx 0.99$ ensures the policy balances immediate throughput against long-term safety.

From MDP to Policy: Q-Values and π

Action-Value Function $Q(s, a)$: The expected discounted return starting from state s , taking action a , and then following the optimal policy forever.

$$Q^*(s, a) = E\{ r + \gamma \max_{a'} Q^*(s', a') \}$$

Intuition:

- $Q(d=0, a=0)$ = high if the queue is empty and communication is valuable
- $Q(d=10, a=0)$ = lower if the queue is full and can't hold more packets
- $Q(w=1, v=1, a=0)$ = very low if weather and speed are unfavorable (high event probability) and we choose communication
- $Q(w=1, v=1, a=1)$ = higher because radar protects us from the high-risk state

Optimal Policy π^* : $\pi^*(s) = \arg \max_{a \in A} Q^*(s, a)$
At each state, pick the action with the highest Q-value.

Challenge: We don't know $Q^*(s, a)$ in advance. We must learn it from experience. That's where Q-learning and DQN come in.

Environment Dynamics: State Transitions $P(s'|s,a)$

Queue Dynamics (Deterministic):

$$d'_{t+1} = \min(d_t - (\text{transmission_count if } a_t=0) + \text{arrivals, } D)$$

Where transmission_count depends on channel state: $v_{\text{max}}=4$ if $c=0$ (good), $v_{\text{max}}=2$ if $c=1$ (bad).
Arrivals $\sim \text{Poisson}(\lambda_d=1)$.

Channel Dynamics (Stochastic, Independent):

$$c'_{t+1} \sim \text{Bernoulli}(p_c=0.1) \rightarrow P(c'=1 | c, a) = 0.1 \text{ (independent of } c \text{ and } a)$$

Environmental Factors (Stochastic, Independent):

Each of r, w, v, m is resampled each step: $P(x'=1) = 1 - \tau_x$. All four factors are independent of each other and of past history (Markov property).

Unexpected Event Generation (Stochastic, Deterministic given state):

$\oplus_t \sim \text{Bernoulli}(P(\oplus))$ where $P(\oplus)$ is computed via Eq. (1). If $a_t=1$ (radar) and $\oplus_t=1$, we detect it. If $a_t=0$ (communication) and $\oplus_t=1$, we miss it (−50 penalty).

Tabular Q-Learning vs. Deep Q-Network (DQN)

Aspect	Tabular Q-Learning	Deep Q-Network (DQN)
Storage	352 table entries (one per state-action pair)	Neural network weights
Update Rule	$Q(s,a) \leftarrow Q(s,a) + \alpha[r + \gamma \max_{a'} Q(s,a') - Q(s,a)]$	Same update rule but applied to network output
Function Approx.	None; exact values	Neural network approximates $Q(s,a)$ for any s
Convergence	~280 episodes (paper results)	170 episodes (paper results)
Scalability	352 states OK; larger problems struggle	Scales to larger state spaces (feature learning)
Training Data Efficiency	One update per experience	Mini-batch updates from replay memory (higher variance reduction)
Exploration	ϵ -greedy (simple)	ϵ -greedy (simple)

Bottom Line: For iRDRC with 352 states, both work. DQN is faster and scales better. Both methods exploit the MDP structure: knowing $P(s'|s,a)$ is unnecessary; learning from experience is enough.

DQN Network Update (Equation 5)

Update Rule:

$$\theta_{t+1} = \theta_t + \alpha [y_t - Q(s_t, a_t; \theta_t)] \nabla Q(s_t, a_t, \theta_t)$$

Where:

- θ_t : Network weights at time t
- α : Learning rate (e.g., 0.001)
- y_t : Target = $r_t + \gamma \max_{a'} Q(s_{t+1}, a'; \theta_t^-)$
- θ_t^- : **Target network** weights (copied from online network periodically, e.g., every 10k steps)
- $\nabla Q(\dots)$: Gradient of Q-value w.r.t. weights

Key Innovation: Target Network

Why use two networks? The target network provides a stable target for learning. Without it, y_t changes as θ changes, causing oscillations. By updating θ^- only periodically, we stabilize training.

DQN Training Loop: 7 Steps per Experience

Step	Action	Details
1	Observe state	AV measures (d, c, r, w, v, m)
2	Estimate Q-values	Feed state to network → get $Q(s, 0)$ and $Q(s, 1)$
3	Choose action	ϵ -greedy: explore (random a) with prob ϵ , exploit (max Q) with prob $1-\epsilon$
4	Interact	Execute action a; environment returns reward r and next state s'
5	Store in memory	Add (s, a, r, s') to replay buffer (circular buffer, capacity ~10k)
6	Sample and train	Random mini-batch (size ~32) from buffer; compute loss and update θ
7	Repeat	Continue for many episodes until cumulative reward converges

Why Replay Memory? Consecutive experiences are highly correlated (e.g., queue decreases only slightly each step). Sampling random mini-batches decorrelates the data, reducing variance and stabilizing learning.

Summary: Complete MDP for iRDRC

Component	Definition	Size/Form	Implementation Note
State Space S	(d, c, r, w, v, m)	352 discrete states	Fully observable; measured sensors
Action Space A	{0: comm, 1: radar}	2 actions	Binary choice each step
Reward R(s,a,s')	$r_{\text{comm}}=2, r_{\text{radar}}=1, r_{\text{coll}}=-50, r_{\text{goal}}=50$ $r_{\text{stop}}=-25$	Scalar	Safety-first asymmetry
Transition P(s' s,a)	Queue (det.); channel, env. (stoch) event prob.	Eq. (10)	Markov: future indep. of history
Discount γ	Typical ≈ 0.99	Scalar in [0,1)	Not specified in paper; choose typical value
Solver	DQN (Keras/PyTorch)	Neural net (6→hidden→17)	170 episodes to converge

Key Takeaways: Why MDP + DQN Works

- 1. MDP Captures the Structure:** The Markov property holds; past doesn't matter, only current state. This allows algorithms to plan without knowing transition probabilities.
- 2. Reward Design Encodes Goals:** The asymmetric reward (-50 ■ $+1, +2$) directly implements a safety priority. Learning doesn't require explicit rules; it discovers the right balance.
- 3. Discount Factor Ensures Long-Horizon Planning:** $\gamma \approx 0.99$ makes the policy care about future crashes. Without it, the policy would be greedy and reckless.
- 4. DQN Scales Beyond Tabular Methods:** 352 states is manageable for both Q-learning and DQN. DQN is faster (170 vs. 280 episodes) due to mini-batch updates and replay memory stabilization.
- 5. Adaptation Emerges from Learning:** The policy is not programmed to 'use radar when risky.' It learns this strategy by optimizing cumulative reward, adapting to any change in p^1 ■ or other conditions.

Modeling Gaps to Address Before Implementation

Gap	What's Missing	Action for Implementation
Probability Values	p_{win} , p^1_{win} , p_{loss} , p^1_{loss} not specified	Choose reasonable estimates (e.g., $p_{\text{win}}=0.01$, $p^1_{\text{win}}=0.05$)
State Evolution	τ_{win} (probability of favorable state) not specified	Assume $\tau_{\text{win}} \approx 0.7$, $\tau_{\text{w}} \approx 0.8$, $\tau_{\text{loss}} \approx 0.6$, $\tau_{\text{m}} \approx 0.7$
Discount Factor γ	Not specified in paper	Use $\gamma \approx 0.99$ (standard choice)
Episode Length T	Not explicit; appears to be 500–1000 steps	Experiment with $T=500$, measure convergence
DQN Hyperparams	Learning rate α , batch size, buffer size not specified	Use standard DQN defaults: $\alpha=0.001$, batch=32, buffer=10k, ϵ -decay linear
Network Architecture	Number of hidden layers and units not specified	Try 2 hidden layers, 64 units each (6→64→64→2)

Ready for Implementation

- ✓ MDP structure fully understood: state, action, reward, transition, objective
- ✓ Why each component matters: Markov property, safety-first reward, discount factor
- ✓ DQN algorithm explained: network update, replay memory, target network
- ✓ Gaps identified and solutions proposed

Next Meeting (Discussion 3): Implementation code walkthrough and first results

Questions & Discussion